

Linking

Introduction to Computer Systems
15th Lecture, Nov. 6, 2025

Instructors:

Class 1: Chen Xiangqun, Liu Xianhua

Class 2: Guan Xuetao

Class 3: Lu Junlin

Outline of Linking

- **Linking: combining object files into programs**
 - Object files
 - Linking mechanism
 - Symbols and symbol resolution
 - Relocation
- **Libraries**
- **Dynamic linking, loading & execution**
- **Library inter-positioning**

Example C Program

```
int sum(int *a, int n);

int array[2] = {1, 2};

int main(int argc, char** argv)
{
    int val = sum(array, 2);
    return val;
}
```

main.c

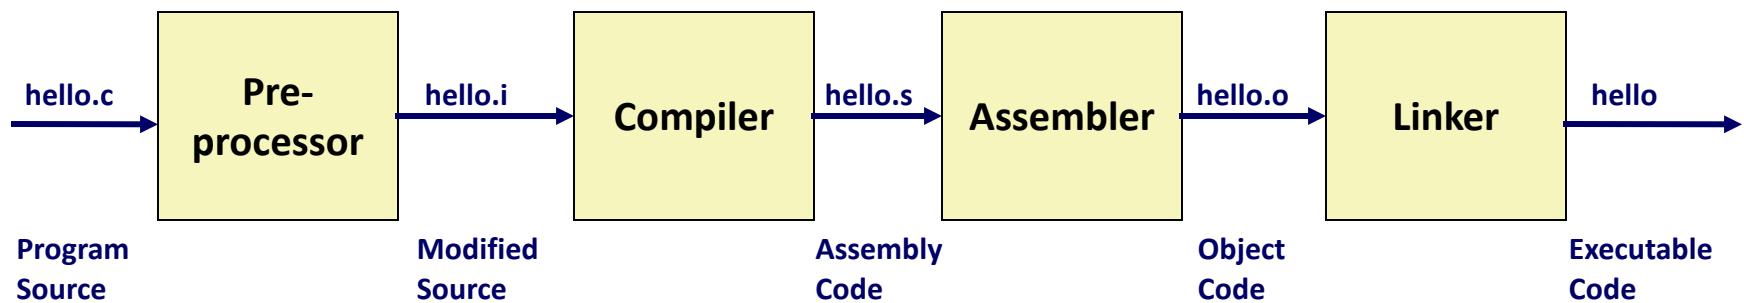
```
int sum(int *a, int n)
{
    int i, s = 0;

    for (i = 0; i < n; i++) {
        s += a[i];
    }
    return s;
}
```

sum.c

Compiler Driver, GCC as an Example

- Gcc is the compiler driver in compilation toolchain.
- Gcc invokes several other compilation phases
 - cpp, the preprocessor
 - cc1, the compiler
 - as/gas, the assembler
 - ld, the linker
- What does each one do? What are their outputs?



Preprocessor

■ First, *gcc* compiler driver invokes *cpp* to generate expanded source

- Preprocessor just does text substitution/ gcc with option “-E”
- Converts the C source file to another C source file
- Expands “#” directives

```
#include <stdio.h>
#define FOO 4
int main() {
    printf("hello, world %d\n", FOO);
}

...
extern int printf (const char * __restrict __format,
      ...);

...
int main() {
    printf("hello, world %d\n", 4);
}
```



Compiler

- Next, *gcc* invokes *cc1* to generate assembly code
 - Translates high-level C code into assembly

```
...  
extern int printf (const char * __restrict __format,  
...);  
...
```

```
int main() {  
    printf("hello, world %d\n", 4);  
}
```

```
.section      .rodata  
.LC0:  
.string "hello, world %d\n"
```

```
.text  
main:  
    pushq  %rbp  
    movq   %rsp, %rbp  
    movl   $4, %esi  
    movl   $.LC0, %edi  
    movl   $0, %eax  
    call   printf  
    popq   %rbp  
    ret
```

Assembler

■ Furthermore, *gcc* invokes *gas* to generate object code

- Translates assembly code into binary object code

```
# readelf -a hello | grep rodata
[10] .rodata           PROGBITS          0000000000495d40  00095d40

# readelf -a hello | grep -E "GLOBAL.* main"
1591: 0000000000401190    31 FUNC    GLOBAL DEFAULT    6 main

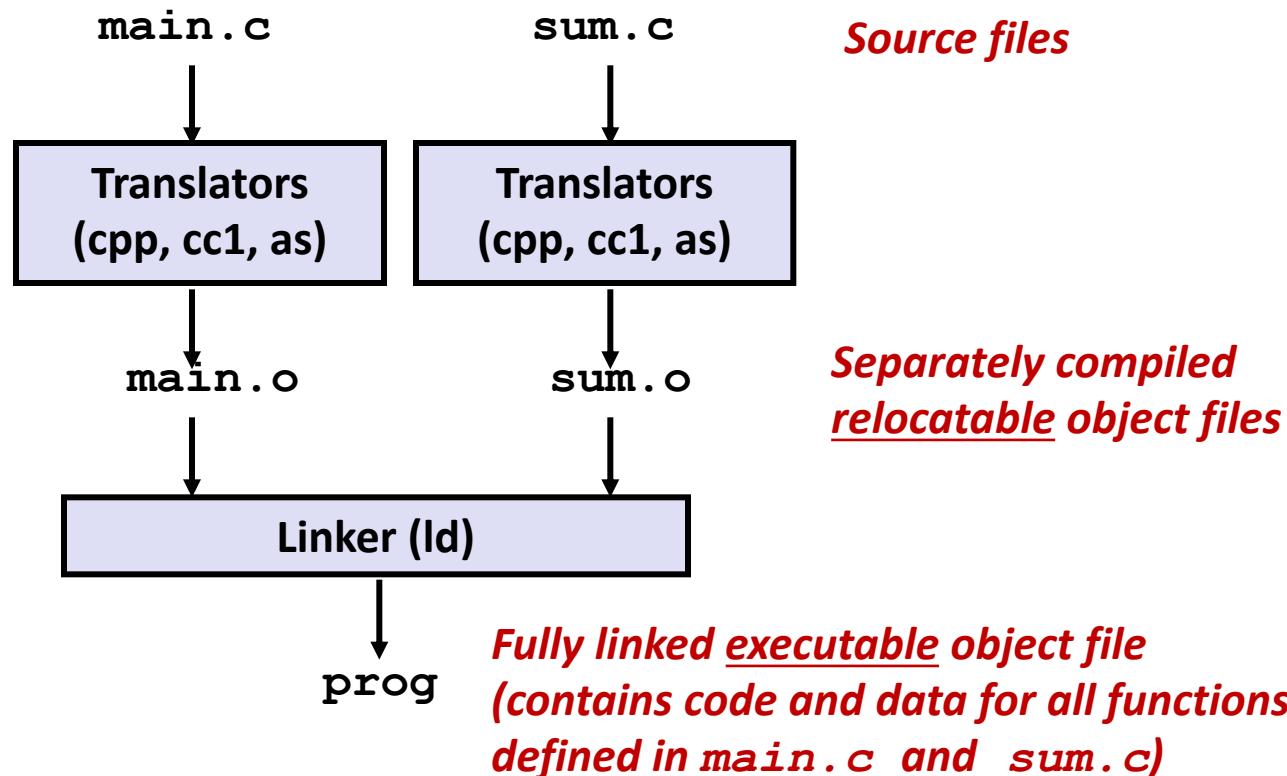
# readelf -x .rodata hello
Hex dump of section '.rodata':
0x00495d40 01000200 68656c6c 6f2c2077 6f726c64 ....hello, world
0x00495d50 2025640a 00464154 414c3a20 6b65726e %d..FATAL: kern

# objdump -d hello
0000000000401190 <main>:
401190:      55                      push   %rbp
401191: 48 89 e5                mov    %rsp,%rbp
401194:  be 04 00 00 00        mov    $0x4,%esi
401199:  bf 44 5d 49 00        mov    $0x495d44,%edi
40119e:  b8 00 00 00 00        mov    $0x0,%eax
4011a3:  e8 d8 0e 00 00        callq 402080 <_IO_printf>
4011a8:  b8 00 00 00 00        mov    $0x0,%eax
4011ad:  5d                  pop    %rbp
4011ae:  c3                  retq
4011af:  90                  nop
```

(Static) Linking

- Programs are translated and linked using a *compiler driver*:

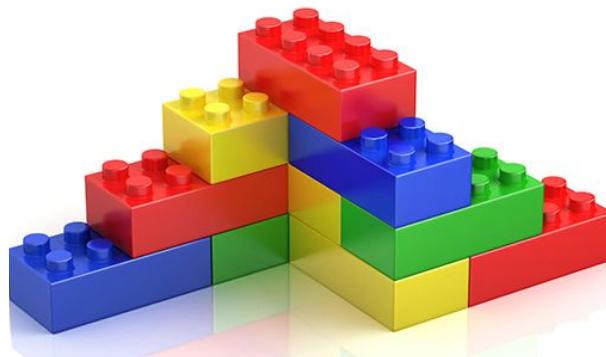
- linux> `gcc -Og -o prog main.c sum.c`
- linux> `./prog`



Why Linkers?

■ Reason 1: Modularity

- Program can be written as a collection of smaller source files, rather than one monolithic mass.
- Can build libraries of common functions (more on this later)
 - e.g., Math library, standard C library



Why Linkers? (cont)

■ Reason 2: Efficiency

- Time: Separate compilation
 - Change one source file, compile, and then relink.
 - No need to recompile other source files.
 - Can compile multiple files concurrently.
- Space: Libraries
 - Common functions can be aggregated into a single file...
 - **Option 1: *Static Linking***
 - Executable files and running memory images contain only the library code they actually use
 - **Option 2: *Dynamic linking***
 - Executable files contain no library code
 - During execution, single copy of library code can be shared across all executing processes

What Do Linkers Do?

■ Step 1: Symbol resolution

- Programs define and reference *symbols* (global variables and functions):
 - `void swap() { ... } /* define symbol swap */`
 - `swap(); /* reference symbol swap */`
 - `int *xp = &x; /* define symbol xp, reference x */`
- Symbol definitions are stored in object file (by assembler) in *symbol table*.
 - Symbol table is an array of entries.
 - Each entry includes name, size, and location of symbol.
- **During symbol resolution step, the linker associates each symbol reference with exactly one symbol definition.**

Symbols in Example C Program

Declaration

```

int sum(int *a, int n);
int sum1(int *a, int n);

int array[2] = {1, 2};

int main(int argc, char** argv)
{
    int val = sum(array, 2);
    return val;
}

```

Definitions

```

int sum(int *a, int n)
{
    int i, s = 0;

    for (i = 0; i < n; i++) {
        s += a[i];
    }
    return s;
}

```

sum.c

```

# gcc -c -o main.o main.c
# gcc -c -o sum.o sum.c
# nm main.o
0000000000000000 D array
0000000000000000 T main
          U sum

# nm sum.o
0000000000000000 T sum

```

You may also try:
objdump -t main.o
objdump -t sum.o

Assembler will not create the entry for function declarations (**sum1** in **main.o**)

What Do Linkers Do? (cont)

■ Step 2: Relocation

- Merges separate code and data sections into single sections
- Relocates symbols from their relative locations in the .o files to their final absolute memory locations in the executable.
- Updates all references to these symbols to reflect their new positions.

Let's look at these two steps in more detail....

Three Kinds of Object Files (Modules)

■ Relocatable object file (.o file)

- Contains code and data in a form that can be combined with other relocatable object files to form executable object file.
 - Each .o file is produced from exactly one source (.c) file

■ Executable object file (a .out file)

- Contains code and data in a form that can be copied directly into memory and then executed.

■ Shared object file (.so file)

- Special type of relocatable object file that can be loaded into memory and linked dynamically, at either load time or run-time.
- Called *Dynamic Link Libraries* (DLLs) by Windows

Executable and Linkable Format (ELF)

- Standard binary format for object files
- One unified format for
 - Relocatable object files (.o),
 - Executable object files (a.out)
 - Shared object files (.so)
- Generic name: ELF binaries
- First appeared in System V Release 4 Unix, c. 1989
- Linux switched to ELF c. 1995, BSD later at c. 1998-2000

ELF Object File Format

■ ELF header

- Word size, byte ordering, file type (.o, exec, .so), machine type, etc.

■ Segment header table

- Page size, virtual address memory segments (sections), segment sizes.

■ .text section

- Code

■ .rodata section

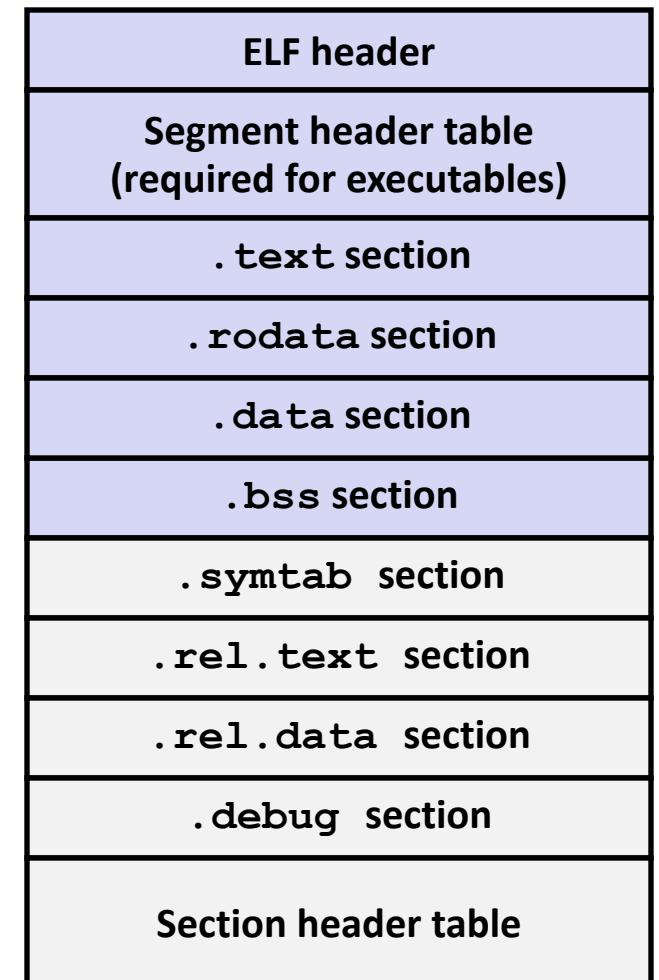
- Read only data: jump tables, string constants, ...

■ .data section

- Initialized global variables

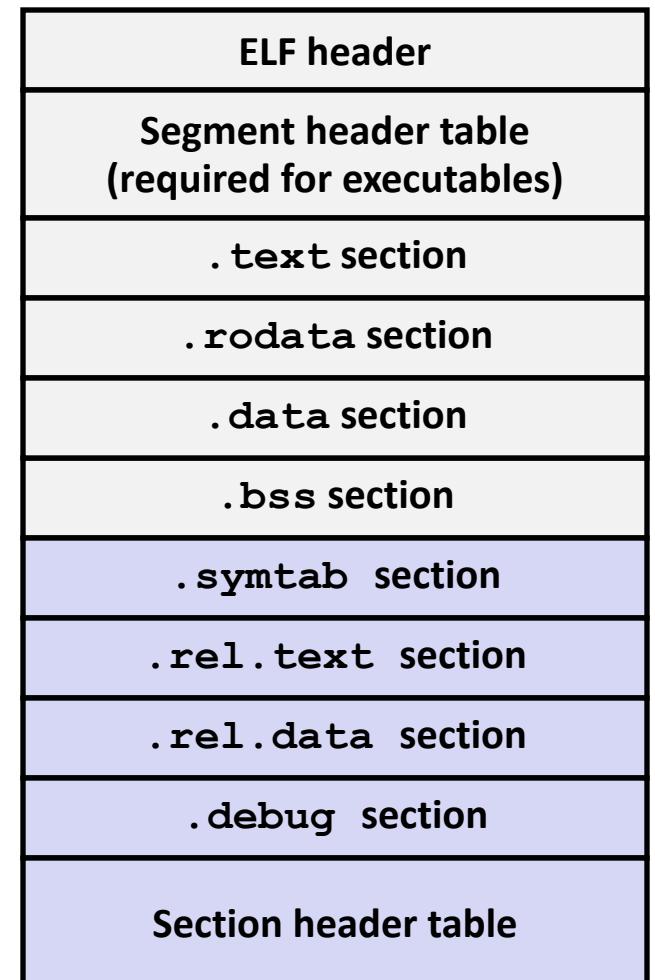
■ .bss section

- Uninitialized global variables
- “Block Started by Symbol”
- “Better Save Space”
- Has section header but occupies no space



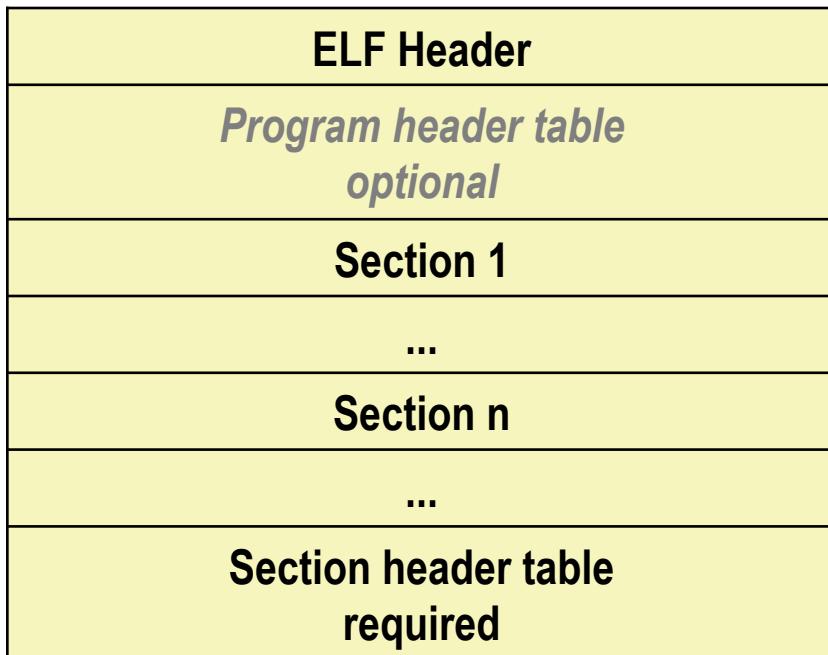
ELF Object File Format (cont.)

- **.syntab section**
 - Symbol table
 - Procedure and static variable names
 - Section names and locations
- **.rel.text section**
 - Relocation info for **.text** section
 - Addresses of instructions that will need to be modified in the executable
 - Instructions for modifying.
- **.rel.data section**
 - Relocation info for **.data** section
 - Addresses of pointer data that will need to be modified in the merged executable
- **.debug section**
 - Info for symbolic debugging (`gcc -g`)
- **Section header table**
 - Offsets and sizes of each section

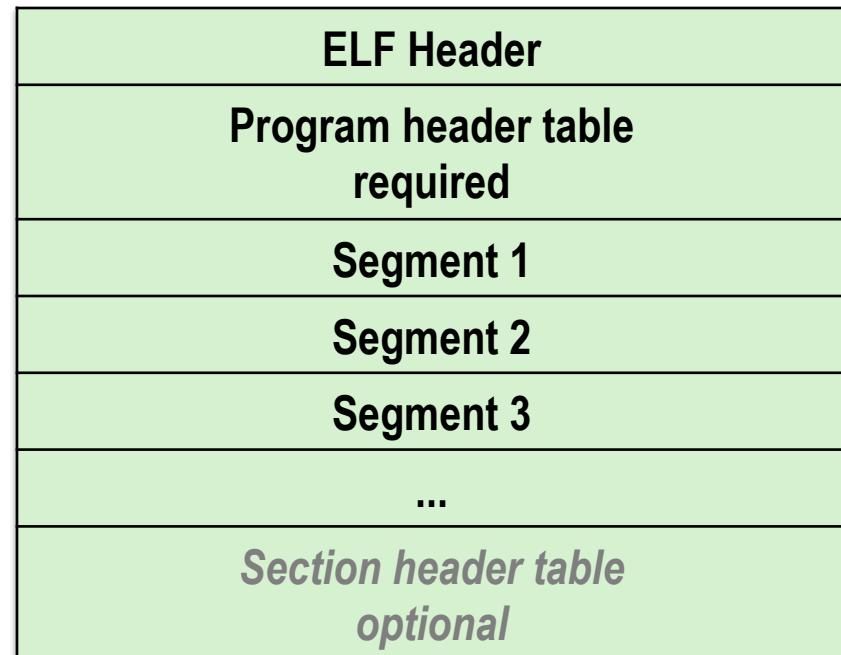


Parallel Views of a ELF File

- ***Program header table/Segments*** is used to build a process image (execute a program); relocatable files don't need it.
- Files used during linking must have a ***section header table/Sections***.



Linking View



Execution View

Linker Symbols

■ Global symbols

- Symbols defined by module m that can be referenced by other modules.
- E.g.: non-**static** C functions and non-**static** global variables.

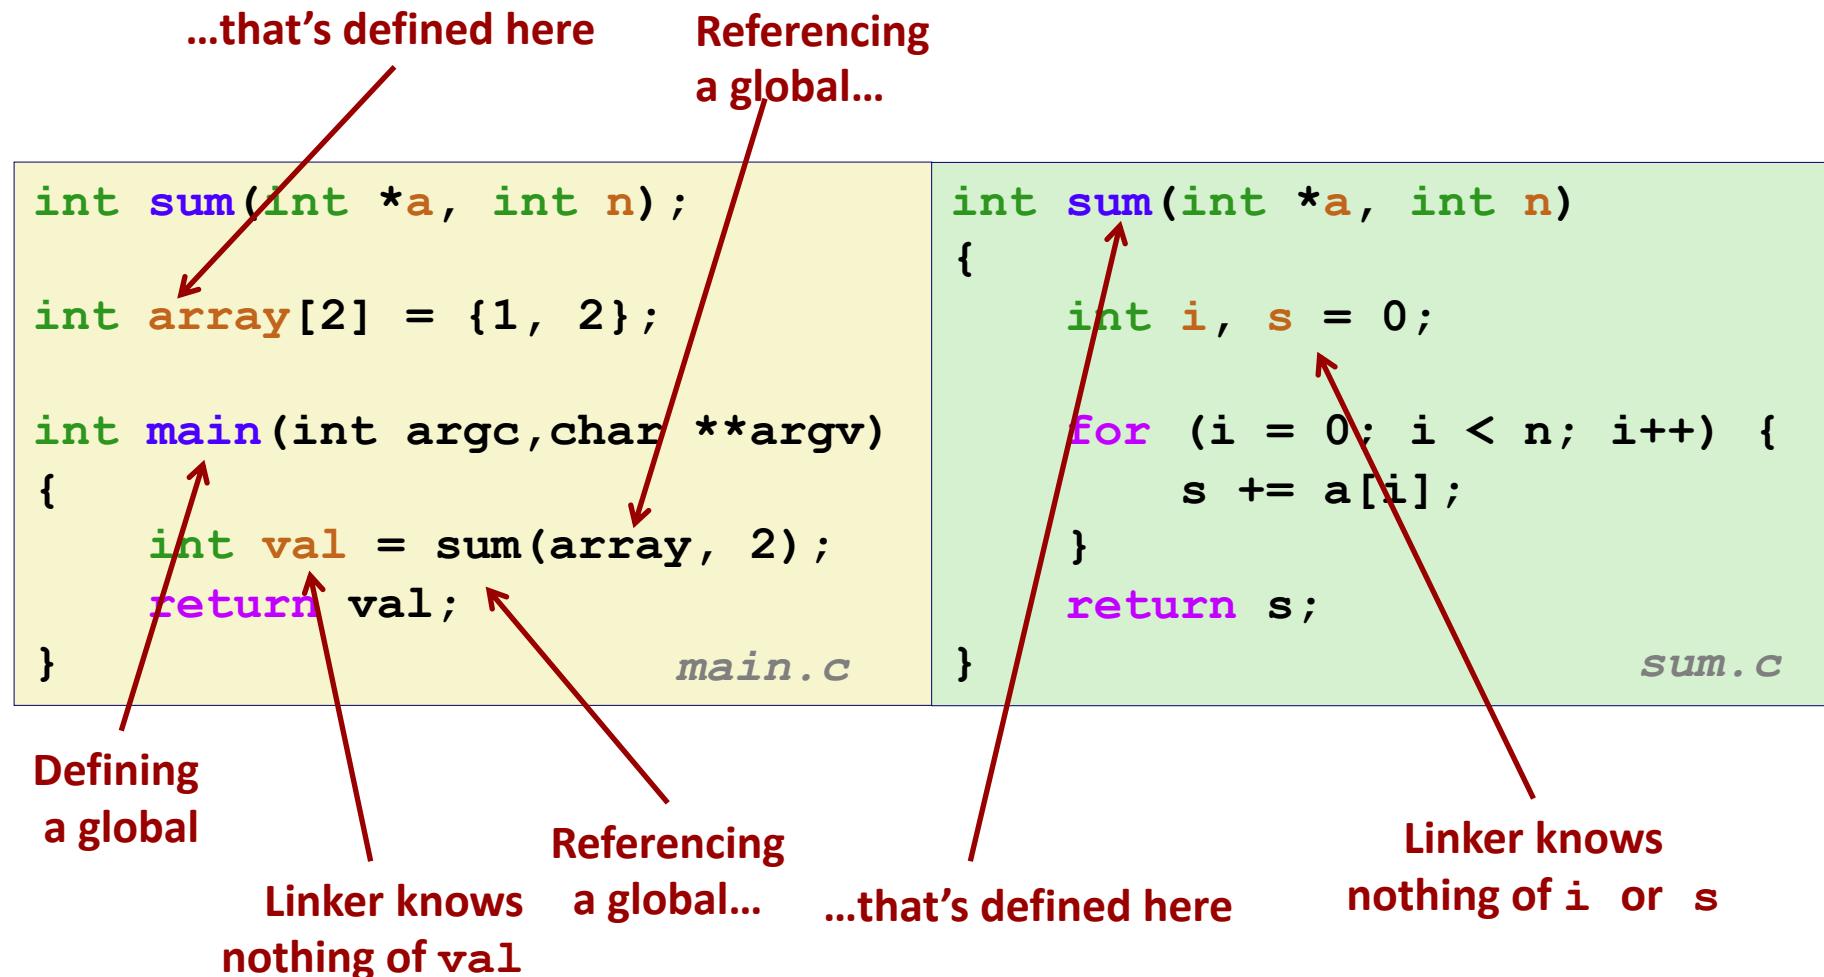
■ External symbols

- Global symbols that are referenced by module m but defined by some other module.

■ Local symbols

- Symbols that are defined and referenced exclusively by module m .
- E.g.: C functions and global variables defined with the **static** attribute.
- **Local linker symbols are *not* local program variables**

Step 1: Symbol Resolution



How Linker Resolves Duplicate Symbol Definitions (such as `sum`, `array`)?

Local Symbols

■ Local non-static C variables vs. local static C variables

- local non-static C variables: stored on the stack
- local static C variables: stored in either .bss, or .data

```
static int x = 15;

int f() {
    static int x = 17;
    return x++;
}

int g() {
    static int x = 19;
    return x += 14;
}

int h() {
    return x += 27;
}
```

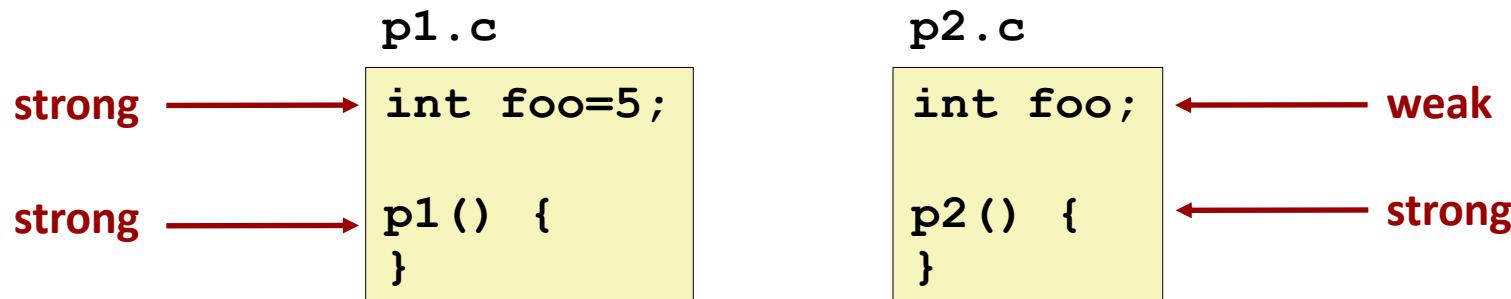
static-local.c

Compiler allocates space in .data for each definition of x

Creates local symbols in the symbol table with unique names, e.g., x, x.1721 and x.1724.

How Linker Resolves Duplicate Symbol Names

- Program symbols are either ***strong*** or ***weak***
 - ***Strong***: procedures and initialized global variables
 - ***Weak***: uninitialized global variables
 - Or ones declared with specifier **`extern`**
- Compiler exports such kind of information and assembler encodes it implicitly in the symbol table of ELF files.



Linker's Symbol Rules

- **Rule 1: Multiple strong symbols are not allowed**
 - Each item can be defined only once
 - Otherwise: Linker error
- **Rule 2: Given a strong symbol and multiple weak symbols, choose the strong symbol**
 - References to the weak symbol resolve to the strong symbol
- **Rule 3: If there are multiple weak symbols, pick an arbitrary one**
 - Can override this with `gcc -fno-common`

Linker Puzzles

```
int x;  
p1() {}
```

```
p1() {}
```

Link time error: two strong symbols (**p1**)

```
int x;  
p1() {}
```

```
int x;  
p2() {}
```

References to **x** will refer to the same uninitialized int. Is this what you really want?

```
int x;  
int y;  
p1() {}
```

```
double x;  
p2() {}
```

Writes to **x** in **p2** might overwrite **y**!
Evil!

```
int x=7;  
int y=5;  
p1() {}
```

```
double x;  
p2() {}
```

Writes to **x** in **p2** might overwrite **y**!
Nasty!

```
int x=7;  
p1() {}
```

```
int x;  
p2() {}
```

References to **x** will refer to the same initialized variable.

Important: Linker does not do type checking.

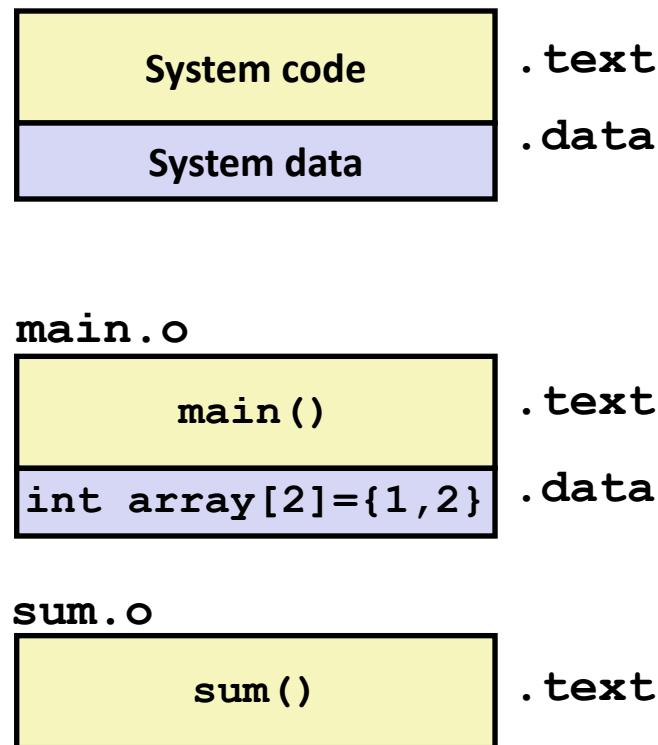
Nightmare scenario: two identical weak structs, compiled by different compilers with different alignment rules.

Rules for avoiding type mismatches

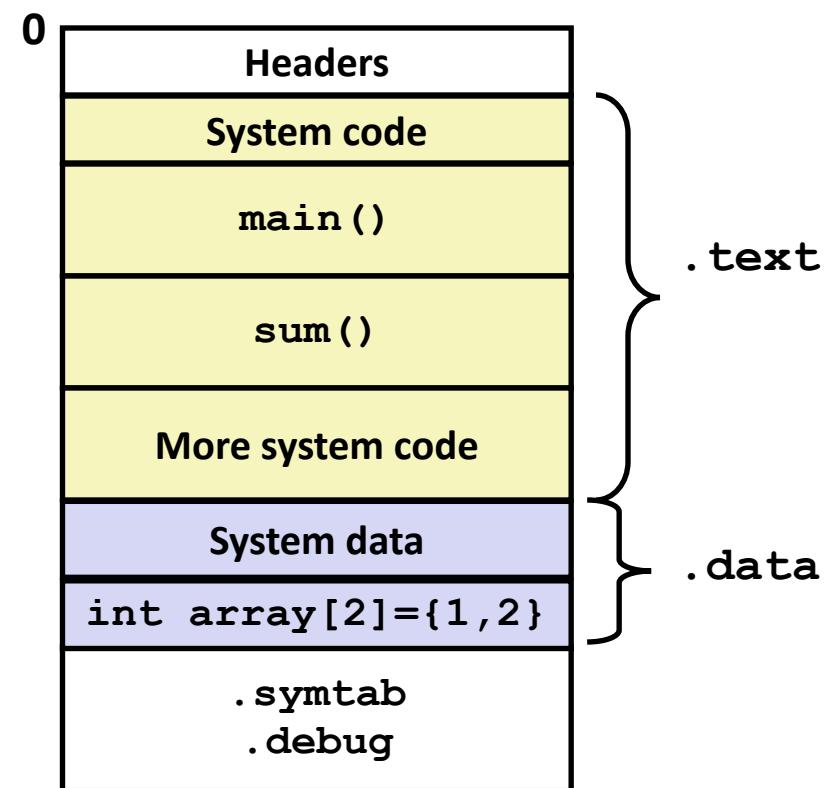
- Avoid global variables as much as possible
- Use **static** as much as possible
- Declare *everything* that's not **static** in a header file
 - Make sure to include the header file everywhere it's relevant
 - Including the files that define those symbols
- Always put **extern** on declarations in header files
 - Unnecessary but harmless for function declarations
 - Avoids the quirky behavior of extern-less global variables
- Always write (**void**) when a function takes no arguments
 - `extern void no_args(void);`
 - Leaving out the **void** means “I'm *not saying* what argument list this function takes.” Turns off argument type checking!

Step 2: Relocation

Relocatable Object Files



Executable Object File



2-Step Relocation in Static Linking

■ Relocating sections and symbol definitions

- Merges all sections of the same type into a new aggregate section of the same type.
- Assigns run-time memory addresses to
 - The new aggregate section.
 - Each section defined by the input modules.
 - Each symbol defined by the input modules.

■ Relocating symbol references with sections

- Modifies every symbol reference in the bodies of the code and data sections so that they point to the correct run-time addresses.
- It relies on data structures in the relocatable modules known as **relocation entries**.

Relocation Entries

- A *relocation entry* generates from reference with unknown location.

```
/* Relocation table entry with addend  
(in section of type SHT_REL). */
```

```
660 typedef struct  
661 {  
662     Elf64_Addr r_offset; /* Address */  
663     Elf64_XWord r_info; /* Relocation type and symbol index */  
664     Elf64_Sxword r_addend; /* Addend */  
665 } Elf64_Rela;  
673 #define ELF64_R_SYM(i) ((i) >> 32)  
674 #define ELF64_R_TYPE(i) ((i) & 0xffffffff)
```

- *r_offset* is the section offset of the reference that will be modified.
- *ELF_64_R_SYM* identifies the symbol that the reference should point to.
- *ELF_64_R_TYPE* tells the linker how to modify the new reference.
- *r_addend* is a constant used for offset adjustment in some kind of relocation.

Two Most Basic Relocation Types

■ R_X86_64_PC32

- Relocates a reference that uses a 32-bit PC-relative address.

■ R_X86_64_32/R_X86_64_32S

- Relocates a reference that uses a 32-bit absolute address.

```
for each section s {
    foreach relocation entry r {
        refptr = s + r.offset; /* ptr to reference to be relocated */

        /* Relocate a PC-relative reference */
        if (r.type == R_X86_64_PC32) {
            refaddr = ADDR(s) + r.offset; /* ref's run-time address */
            *refptr = (unsigned) (ADDR(r.symbol) + r.addend - refaddr);
        }

        /* Relocate an absolute reference */
        if (r.type == R_X86_64_32)
            *refptr = (unsigned) (ADDR(r.symbol) + r.addend);
    }
}
```

Relocation Entries

```

int array[2] = {1, 2};
int main(int argc, char** argv) {
    int val = sum(array, 2);
    return val;
}

```

main.c

# readelf -r main.o			
Relocation section '.rela.text' at offset 0x560 contains 2 entries:			
Offset	Info	Type	Sym.Value
000000000015	00080000000a	R_X86_64_32	0000000000000000
00000000001a	000a00000002	R_X86_64_PC32	0000000000000000
Relocation section '.rela.eh_frame' at offset 0x590 contains 1 entries:			
Offset	Info	Type	Sym.Name+Addend
000000000020	000200000002	R_X86_64_PC32	.text + 0
offset		type	symbol name & addend

Totally 3 symbols to be relocated.

Relocation Entries (in main.o)

```

int array[2] = {1, 2};
int main(int argc, char** argv) {
    int val = sum(array, 2);
    return val;
}

```

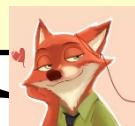
main.c

```
# readelf -r main.o
Relocation section '.rela.text' at offset 0x560 contains 2 entries:
Offset      Info      Type            Sym.Value      Sym.Name+Addend
000000000015 00080000000a R_X86_64_32 0000000000000000 array + 0
00000000001a 000a00000002 R_X86_64_PC32 0000000000000000 sum - 4
```

Dear Linker,

Please patch the *.rela.text* section at offsets 0x15. Patch in a 32-bit value like following steps. When you determine the addr of *.data*, compute [addr of array] + [addend, which equals 0] and place the result at the prescribed place.

Sincerely,
Assembler



Relocation Entries (in main.o)

```

int array[2] = {1, 2};
int main(int argc, char** argv) {
    int val = sum(array, 2);
    return val;
}

```

main.c

```

# readelf -r main.o
Relocation section '.rela.text' at offset 0x560 contains 2 entries:
Offset      Info      Type            Sym.Value      Sym.Name+Addend
000000000015 00080000000a R_X86_64_32    0000000000000000 array + 0
00000000001a 000a00000002 R_X86_64_PC32  0000000000000000 sum - 4

```

Dear Linker,

Please patch the .rela.text section at offsets 0x1a. Patch in a 32-bit “PC-relative” value like following steps. When you determine the addr of sum, compute [addr of sum] + [addend, which equals -4] – [addr of section + offset] and place the result at the prescribed place.

Sincerely,
Assembler



Relocation Entries (in sum.o)

```

int sum(int *a, int n)
{
    int i, s = 0;
    for (i = 0; i < n; i++) {
        s += a[i];
    }
    return s;
}

```

sum.c

# readelf -r sum.o	Relocation section '.rela eh_frame' at offset 0x4f8 contains 1 entries:	Sym. Name+Addend
Offset	Type	Sym. Value
000000000020	R_X86_64_PC32	0000000000000000

offset

type

symbol name & addend

1 symbol to be relocated (.text)

Original Object File of main.o

```

int array[2] = {1, 2};
int main(int argc, char** argv) {
    int val = sum(array, 2);
    return val;
}

```

main.c

<pre>0000000000000000 <main>:</pre> <pre> 0: 55 push %rbp 1: 48 89 e5 mov %rsp,%rbp 4: 48 83 ec 20 sub \$0x20,%rsp 8: 89 7d ec mov %edi,-0x14(%rbp) b: 48 89 75 e0 mov %rsi,-0x20(%rbp) f: be 02 00 00 00 mov \$0x2,%esi 14: bf 00 00 00 00 mov \$0x0,%edi # %edi = &array 15: R_X86_64_32 array # Relocation entry 19: e8 00 00 00 00 callq 1e <main+0x1e> # sum() 1a: R_X86_64_PC32 sum-0x4 # Relocation entry 1e: 89 45 fc mov %eax,-0x4(%rbp) 21: 8b 45 fc mov -0x4(%rbp),%eax 24: c9 leaveq 25: c3 retq </pre>	<pre>Source: objdump -r -d main.o</pre>
---	---

Original Object File of sum.o

```
int sum(int *a, int n) {
```

0000000000000000 <sum>:

0:	55		push	%rbp
1:	48 89 e5		mov	%rsp,%rbp
4:	48 89 7d e8		mov	%rdi,-0x18(%rbp)
8:	89 75 e4		mov	%esi,-0x1c(%rbp)
b:	c7 45 fc 00 00 00 00		movl	\$0x0,-0x4(%rbp)
12:	c7 45 f8 00 00 00 00		movl	\$0x0,-0x8(%rbp)
19:	eb 1d		jmp	38 <sum+0x38>
1b:	8b 45 f8		mov	-0x8(%rbp),%eax
1e:	48 98		cltq	
20:	48 8d 14 85 00 00 00		lea	0x0(%rax,4),%rdx
27:	00			
28:	48 8b 45 e8		mov	-0x18(%rbp),%rax
2c:	48 01 d0		add	%rdx,%rax
2f:	8b 00		mov	(%rax),%eax
31:	01 45 fc		add	%eax,-0x4(%rbp)
34:	83 45 f8 01		addl	\$0x1,-0x8(%rbp)
38:	8b 45 f8		mov	-0x8(%rbp),%eax
3b:	3b 45 e4		cmp	-0x1c(%rbp),%eax
3e:	7c db		jl	1b <sum+0x1b>
40:	8b 45 fc		mov	-0x4(%rbp),%eax
43:	5d		pop	%rbp
44:	c3		retq	

.text=0xbabe00

```
0000000000000000 <main>:
```

```
 0: 55
 1: 48 89 e5
 4: 48 83 ec 20
 8: 89 7d ec
 b: 48 89 75 e0
 f: be 02 00 00 00
14: bf 00 00 00 00
19: e8 00 00 00 00
1e: 89 45 fc
21: 8b 45 fc
24: c9
25: c3
```

main.o

```
0000000000000000 <sum>:
```

```
 0: 55
 1: 48 89 e5
 4: 48 89 7d e8
 8: 89 75 e4
 b: c7 45 fc 00 00 00 00
12: c7 45 f8 00 00 00 00
19: eb 1d
1b: 8b 45 f8
1e: 48 98
20: 48 8d 14 85 00 00 00
27: 00
28: 48 8b 45 e8
2c: 48 01 d0
2f: 8b 00
31: 01 45 fc
34: 83 45 f8 01
38: 8b 45 f8
3b: 3b 45 e4
3e: 7c db
40: 8b 45 fc
43: 5d
44: c3
```

sum.o

 0000000000000000 <_start>:

```
...
0000000000babf18 <main>:
```

```
babf18: 55
babf19: 48 89 e5
babf1c: 48 83 ec 20
babf20: 89 7d ec
babf23: 48 89 75 e0
babf27: be 02 00 00 00
babf2c: bf 10 fe ea 00
babf31: e8 0a 00 00 00
```

```
babf36: 89 45 fc
babf39: 8b 45 fc
babf3c: c9
babf3d: c3
```

0000000000babf40 <sum>:

```
babf40: 55
babf41: 48 89 e5
babf44: 48 89 7d e8
babf48: 89 75 e4
babf4b: c7 45 fc 00 00 00 00
babf52: c7 45 f8 00 00 00 00
babf59: eb 1d
babf5b: 8b 45 f8
babf5e: 48 98
babf60: 48 8d 14 85 00 00 00
babf67: 00
babf68: 48 8b 45 e8
babf6c: 48 01 d0
babf6f: 8b 00
babf71: 01 45 fc
babf74: 83 45 f8 01
babf78: 8b 45 f8
babf7b: 3b 45 e4
babf7e: 7c db
babf80: 8b 45 fc
babf83: 5d
babf84: c3
```

executable

Disassembly of section .data:

```
0000000000cafe00 <__data_start>:
...
0000000000cafe10 <array>:
cafe10: 00 00
cafe12: 00 00
cafe14: 02 00
```

.data=0xcafe00

**addr_of_array=0xafe10
Using the value**

**0xafe10 to modify the
content here**

addr_of_main=0xbabf18

addr_of_sum=0xbabf40

offset = 0x1a

addend = -4

refptr

= 0xbabf18 + 0x1a

= 0xbabf32

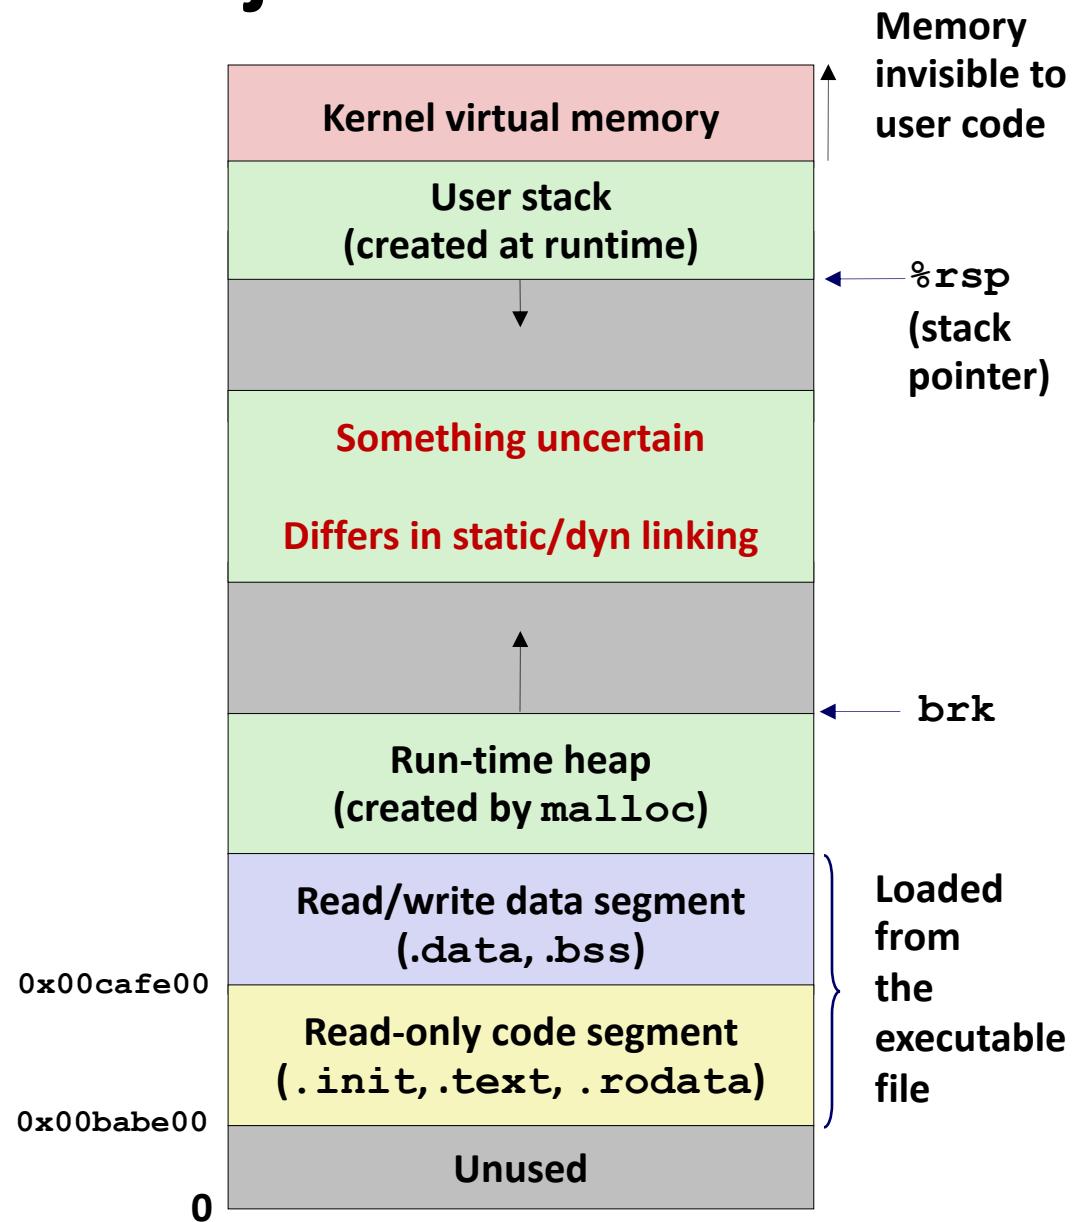
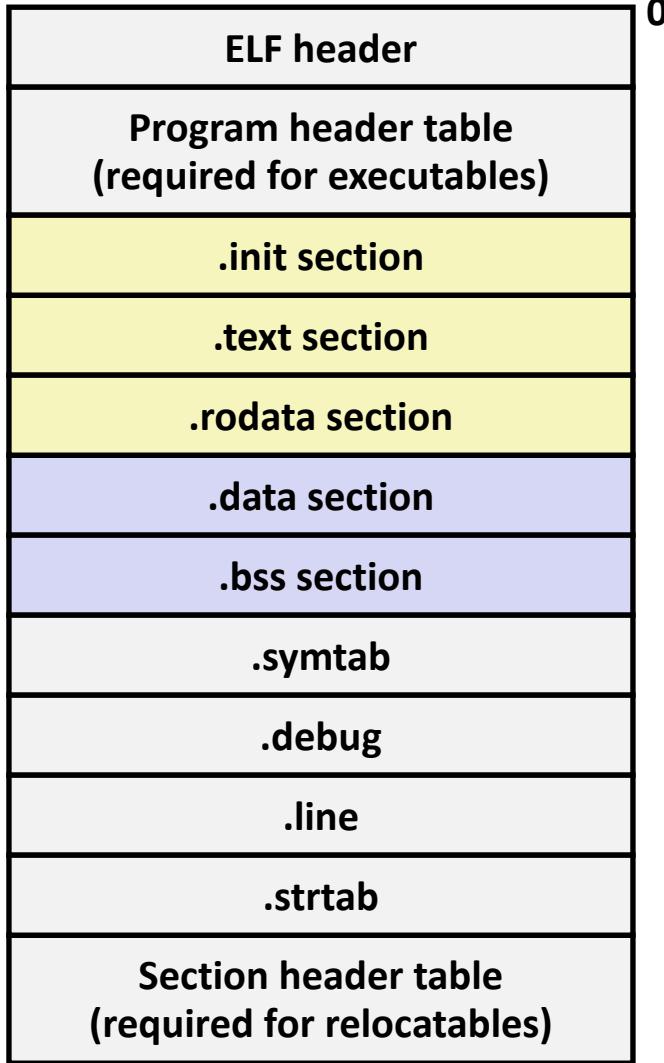
***refptr(content)**

=0xbabf40-4-0xbabf32

=0x0a

Loading Executable Object Files

Executable Object File



Libraries: Packaging a Set of Functions

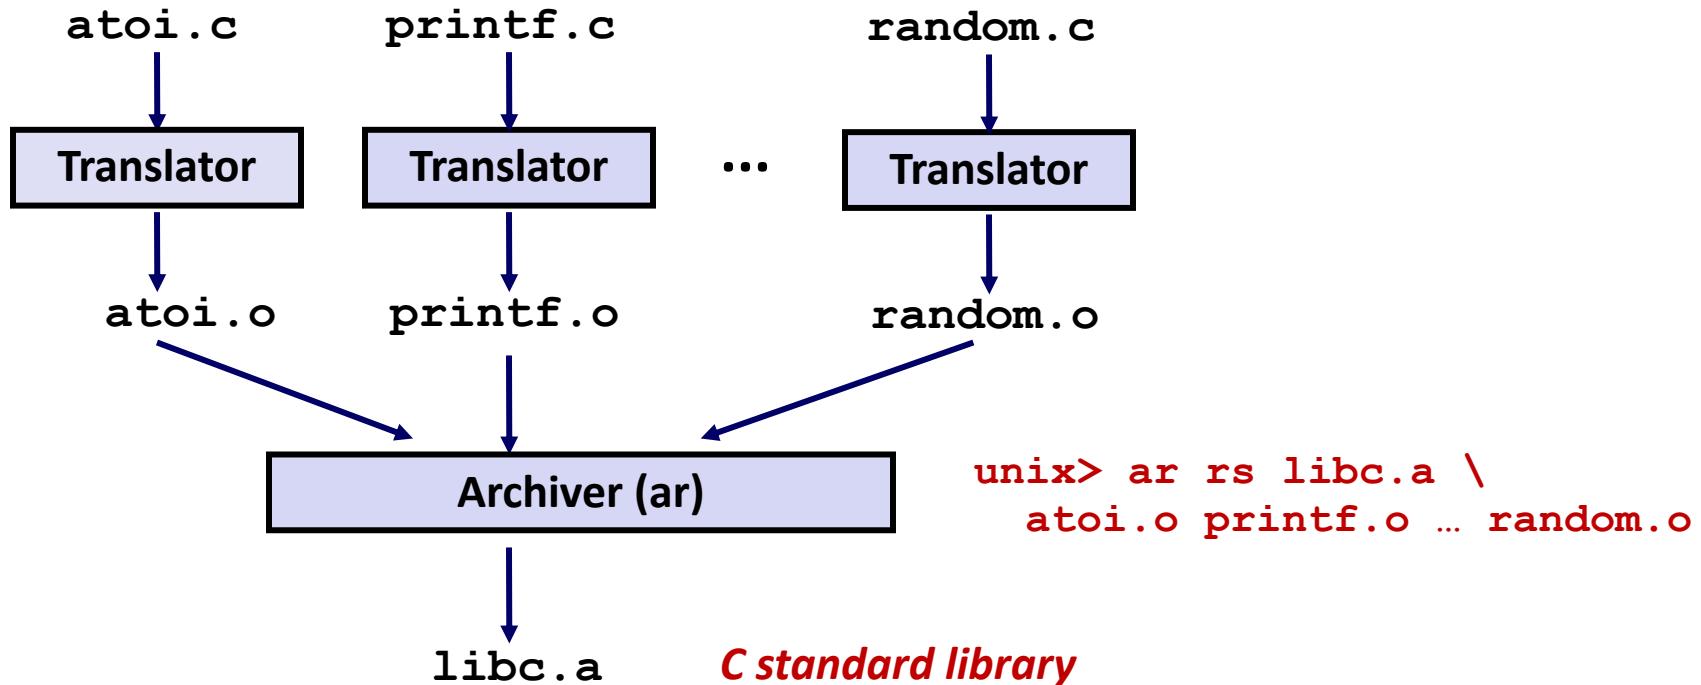
- **How to package functions commonly used by programmers?**
 - Math, I/O, memory management, string manipulation, etc.
- **Awkward, given the linker framework so far:**
 - **Option 1:** Put all functions into a single source file
 - Programmers link big object file into their programs
 - Space and time inefficient
 - **Option 2:** Put each function in a separate source file
 - Programmers explicitly link appropriate binaries into their programs
 - More efficient, but burdensome on the programmer

Old-fashioned Solution: Static Libraries

■ **Static libraries (.a archive files)**

- Concatenate related relocatable object files into a single file with an index (called an *archive*).
- Enhance linker so that it tries to resolve unresolved external references by looking for the symbols in one or more archives.
- If an archive member file resolves reference, link it into the executable.

Creating Static Libraries



- Archiver allows incremental updates
- Recompile function that changes and replace .o file in archive.

Commonly Used Libraries

`libc.a` (the C standard library)

- 4.6 MB archive of 1496 object files. (differs in different versions)
- I/O, memory allocation, signal handling, string handling, data and time, random numbers, integer math

`libm.a` (the C math library)

- 2 MB archive of 444 object files. (differs in different versions)
- floating point math (sin, cos, tan, log, exp, sqrt, ...)

```
% ar -t /usr/lib/libc.a | sort
...
fork.o
...
fprintf.o
fpu_control.o
fputc.o
freopen.o
fscanf.o
fseek.o
fstab.o
...
```

```
% ar -t /usr/lib/libm.a | sort
...
e_acos.o
e_acosf.o
e_acosh.o
e_acoshf.o
e_acoshl.o
e_acosl.o
e_asin.o
e_asinf.o
e_asinl.o
...
```

Linking with Static Libraries

```
#include <stdio.h>
#include "vector.h"

int x[2] = {1, 2};
int y[2] = {3, 4};
int z[2];

int main(int argc, char** argv)
{
    addvec(x, y, z, 2);
    printf("z = [%d %d]\n",
           z[0], z[1]);
    return 0;          main2.c
}
```

libvector.a

```
void addvec(int *x, int *y,
            int *z, int n) {
    int i;

    for (i = 0; i < n; i++)
        z[i] = x[i] + y[i];
}
```

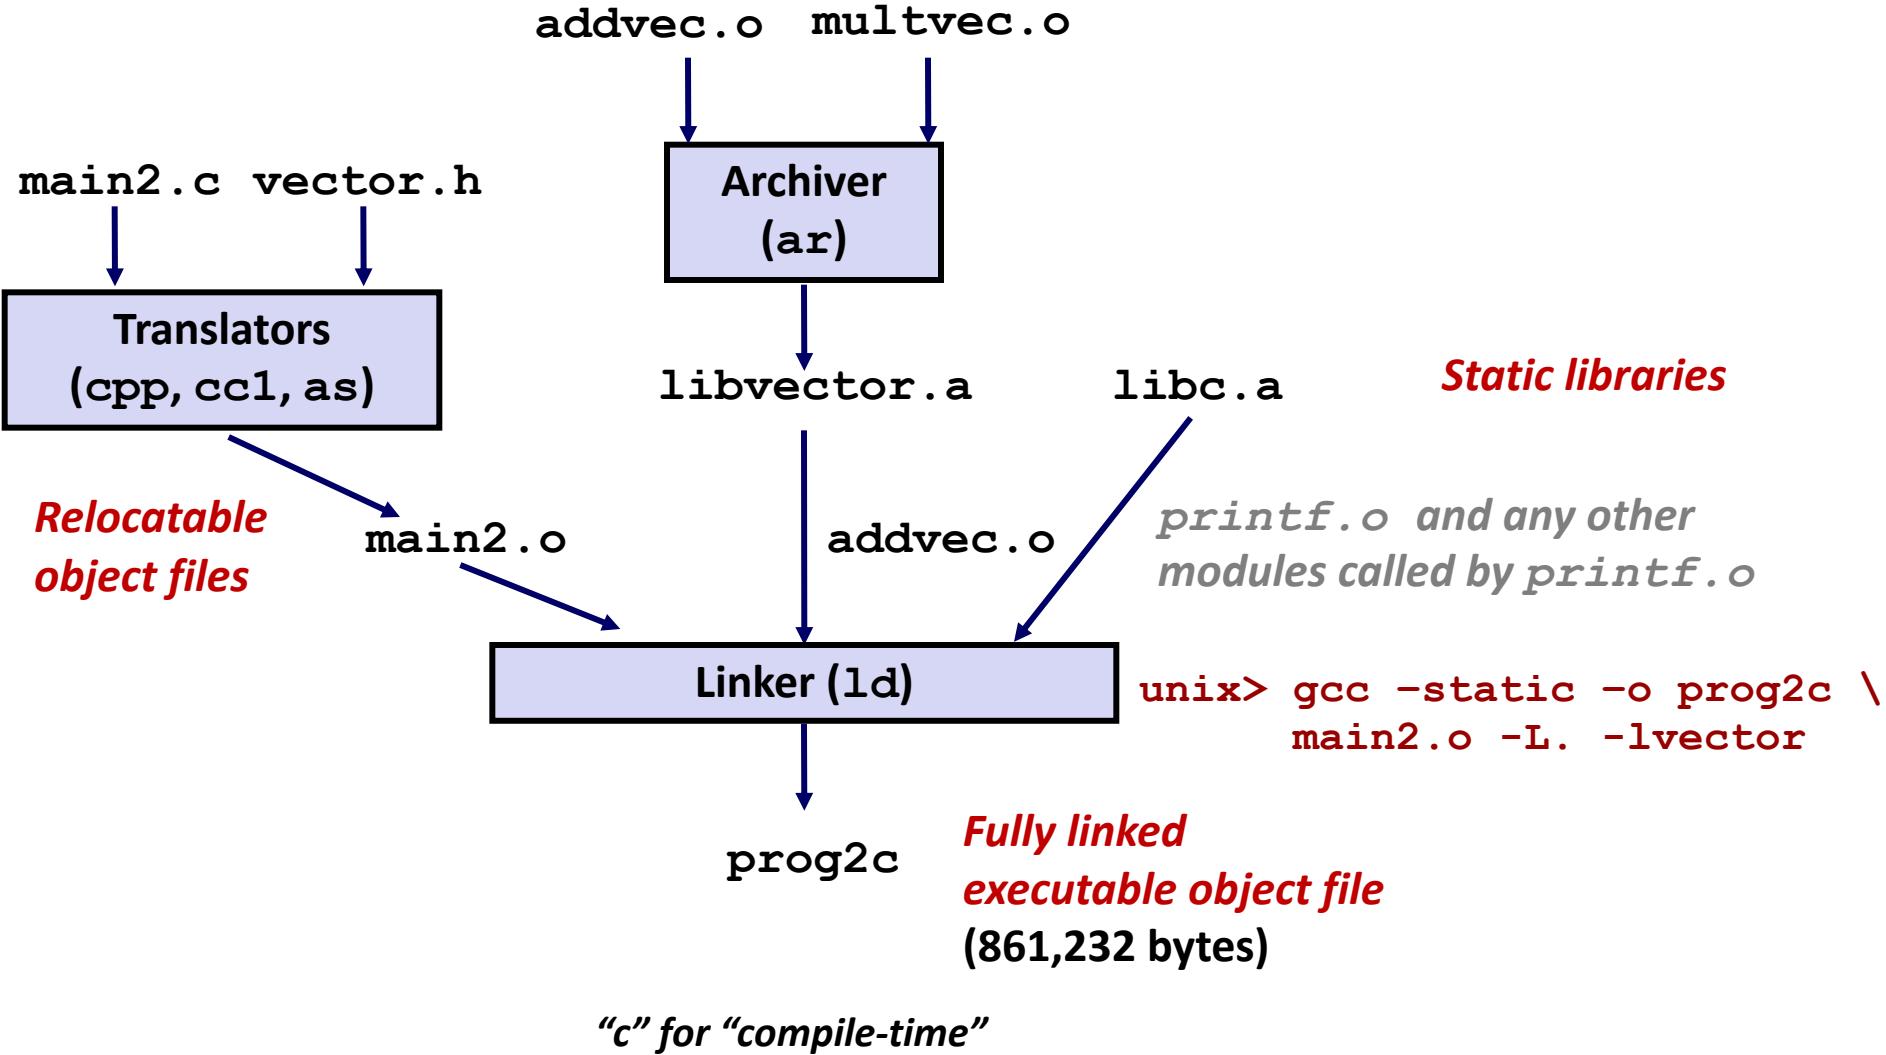
addvec.c

```
void multvec(int *x, int *y,
             int *z, int n)
{
    int i;

    for (i = 0; i < n; i++)
        z[i] = x[i] * y[i];
}
```

multvec.c

Linking with Static Libraries



Using Static Libraries

■ Linker's algorithm for resolving external references:

- Scan **.o** files and **.a** files in the command line order.
- During the scan, keep a list of the current unresolved references.
- As each new **.o** or **.a** file, *obj*, is encountered, try to resolve each unresolved reference in the list against the symbols defined in *obj*.
- If any entries in the unresolved list at end of scan, then error.

■ Problem:

- Command line order matters!
- Moral: put libraries at the end of the command line.

```
unix> gcc -static -o prog2c -L. -lvector main2.o
main2.o: In function `main':
main2.c:(.text+0x19): undefined reference to `addvec'
collect2: error: ld returned 1 exit status
```

Modern Solution: Shared Libraries

■ Static libraries have the following disadvantages:

- Duplication in the stored executables (every function needs libc)
- Duplication in the running executables
- Minor bug fixes of system libraries require each application to explicitly relink
 - Rebuild everything with glibc?
 - <https://security.googleblog.com/2016/02/cve-2015-7547-glibc-getaddrinfo-stack.html>

■ Modern solution: Shared Libraries

- Object files that contain code and data that are loaded and linked into an application *dynamically*, at either *load-time* or *run-time*
- Also called: dynamic link libraries, DLLs, .so files

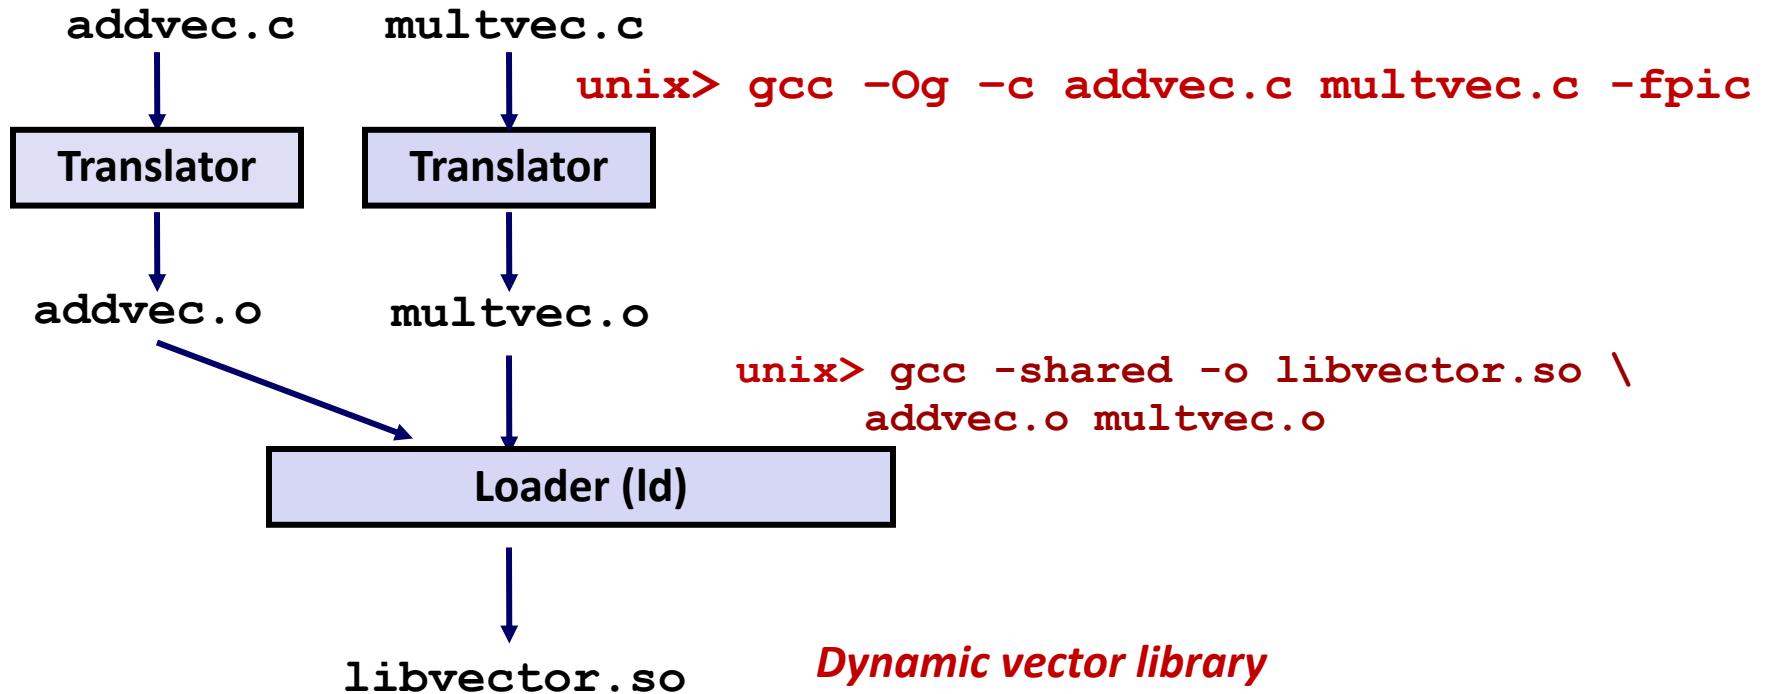
Shared Libraries (cont.)

- Dynamic linking can occur when executable is first loaded and run (load-time linking).
 - Common case for Linux, handled automatically by the dynamic linker (**ld-linux.so**) .
 - Standard C library (**libc.so**) usually dynamically linked.
- Dynamic linking can also occur after program has begun (run-time linking).
 - In Linux, this is done by calls to the **dlopen()** interface .
 - Distributing software.
 - High-performance web servers.
 - Runtime library interpositioning.
- Shared library routines can be shared by multiple processes.
 - More on this when we learn about virtual memory

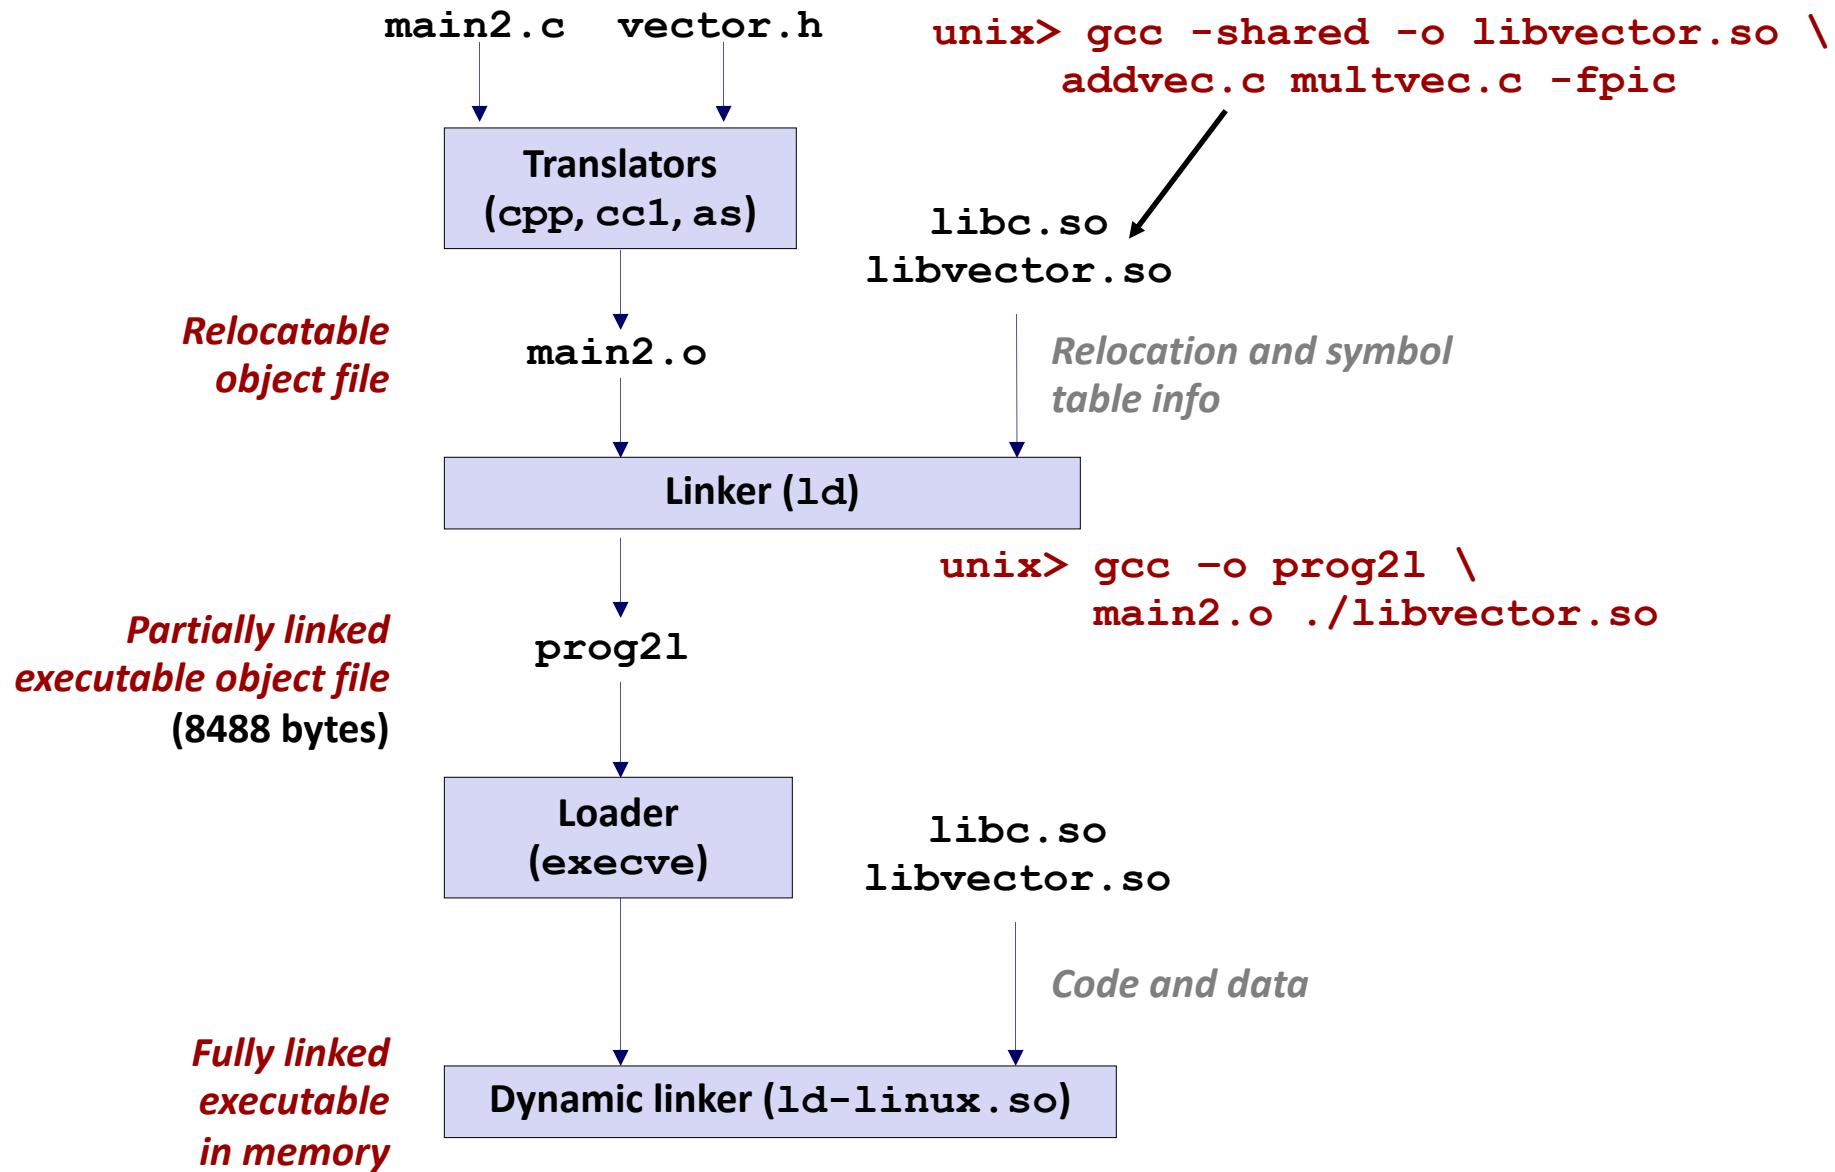
*There..... is... no...hur...ry.
We...love...lazy...binding...*



Dynamic Library Example



Dynamic Linking at Load-time



Dynamic Linking at Run-time

```
#include <stdio.h>
#include <stdlib.h>
#include <dlfcn.h>

int x[2] = {1, 2};
int y[2] = {3, 4};
int z[2];

int main(int argc, char** argv)
{
    void *handle;
    void (*addvec)(int *, int *, int *, int);
    char *error;

    /* Dynamically load the shared library that contains addvec() */
    handle = dlopen("./libvector.so", RTLD_LAZY);
    if (!handle) {
        fprintf(stderr, "%s\n", dlerror());
        exit(1);
    }
    . . .
    addvec(z, x, y, 4);
    printf("z = %d\n", z[1]);
}
```

dll.c

Dynamic Linking at Run-time (cont)

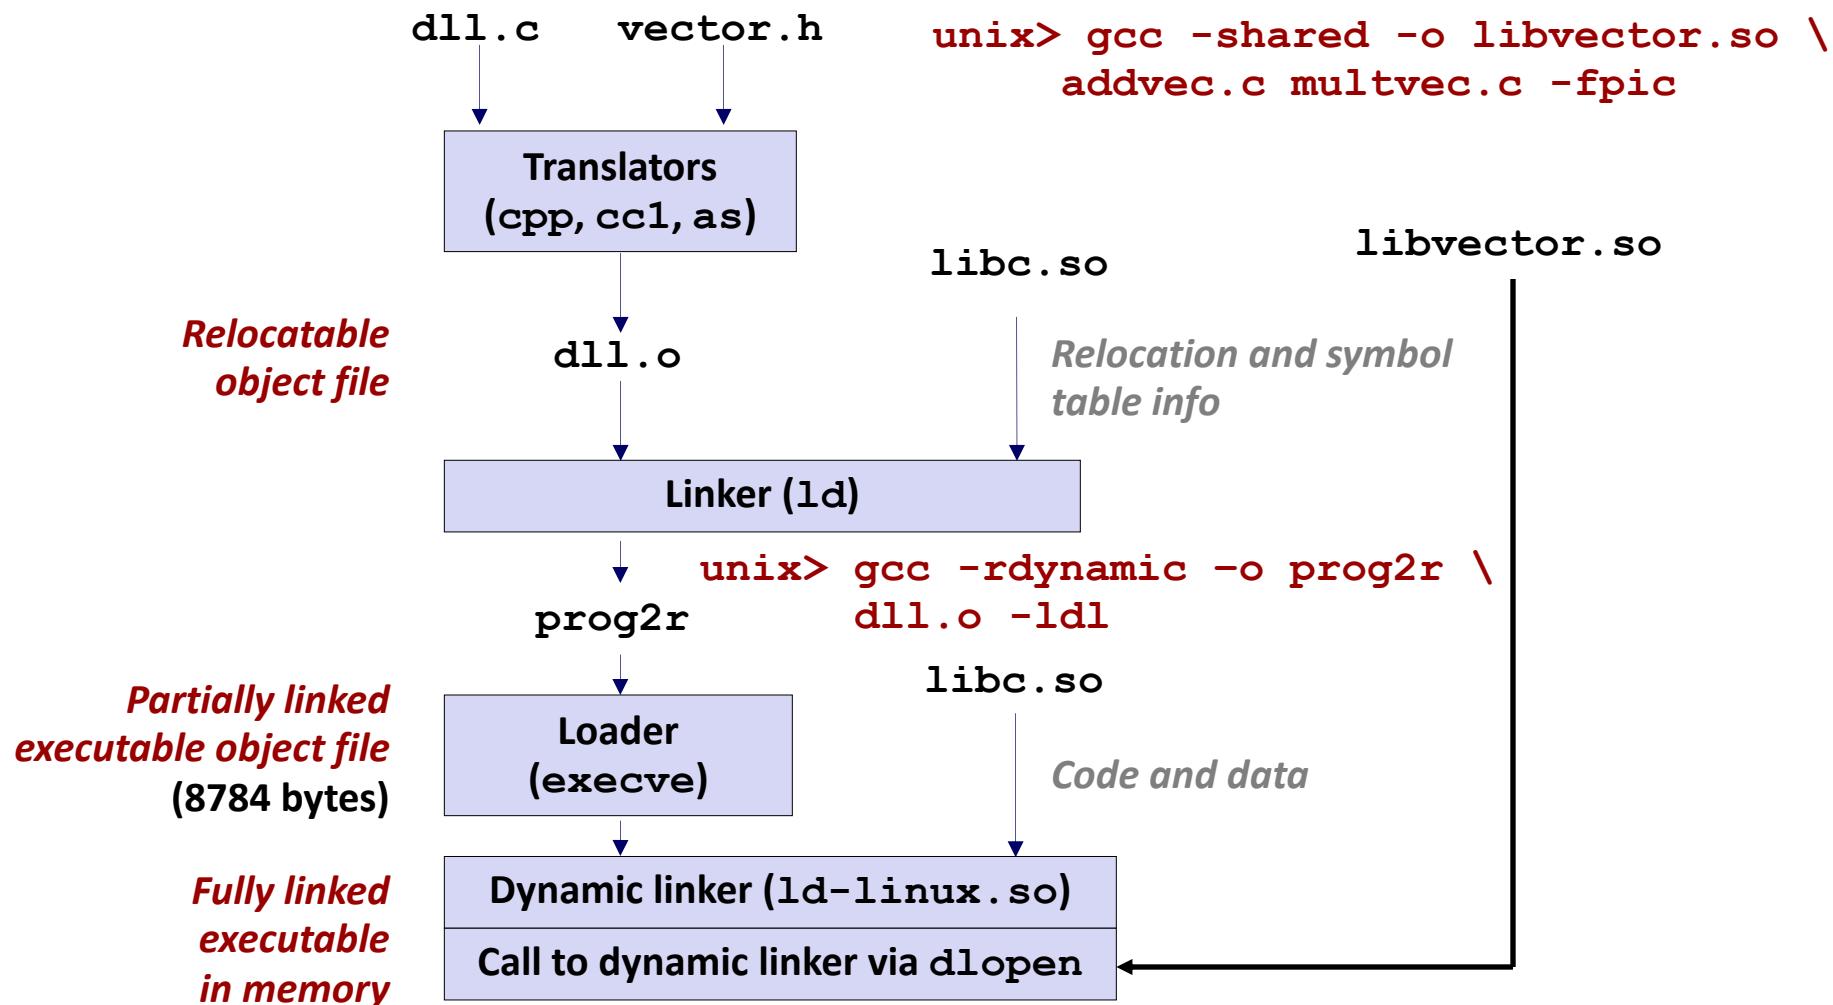
```
...
/* Get a pointer to the addvec() function we just loaded */
addvec = dlsym(handle, "addvec");
if ((error = dlerror()) != NULL) {
    fprintf(stderr, "%s\n", error);
    exit(1);
}

/* Now we can call addvec() just like any other function */
addvec(x, y, z, 2);
printf("z = [%d %d]\n", z[0], z[1]);

/* Unload the shared library */
if (dlclose(handle) < 0) {
    fprintf(stderr, "%s\n", dlerror());
    exit(1);
}
return 0;
}
```

dll.c

Dynamic Linking at Run-time



What dynamic libraries are required?

■ .interp section

- Specifies the dynamic linker to use (i.e., `ld-linux.so`)

■ .dynamic section

- Specifies the names, etc of the dynamic libraries to use
- Follow an example of `prog`

(NEEDED)

Shared library: [libm.so.6]

■ Where are the libraries found?

- Use “`ldd`” to find out:

```
unix> ldd prog
linux-vdso.so.1 => (0x00007ffcf2998000)
libc.so.6 => /lib/x86_64-linux-gnu/libc.so.6 (0x00007f99ad927000)
/lib64/ld-linux-x86-64.so.2 (0x00007f99adcef000)
```

Static vs. Dynamic Linking Tradeoffs

Static:

- Does not need to look up libraries at runtime
- Does not need extra PLT indirection
- Consumes more memory with copies of each library in every program

Dynamic:

- Less disk space/memory (7K vs 571K for hello world)
- Shared libraries already in memory and in hot cache
- Incurs lookup and indirection overheads

Linking Summary

- **Linking is a technique that allows programs to be constructed from multiple object files.**
- **Linking can happen at different times in a program's lifetime:**
 - Compile time (when a program is compiled)
 - Load time (when a program is loaded into memory)
 - Run time (while a program is executing)
- **Understanding linking can help you avoid nasty errors and make you a better programmer.**

Case Study: Library Interpositioning

- Documented in Section 7.13 of textbook
- Library interpositioning : powerful linking technique that allows programmers to intercept calls to arbitrary functions
- Interpositioning can occur at:
 - Compile time: When the source code is compiled.
 - Link time: When the relocatable object files are statically linked to form an executable object file.
 - Load/run time: When an executable object file is loaded into memory, dynamically linked, and then executed.

Some Interpositioning Applications

■ Security

- Confinement (sandboxing)
- Behind the scenes encryption

■ Debugging

- In 2014, two Facebook engineers debugged a treacherous 1-year old bug in their iPhone app using interpositioning
- Code in the SPDY networking stack was writing to the wrong location
- Solved by intercepting calls to POSIX write functions (write, writev, pwrite)

Source: Facebook engineering blog post at:

<https://code.facebook.com/posts/313033472212144/debugging-file-corruption-on-ios/>

Some Interpositioning Applications

■ Monitoring and Profiling

- Count number of calls to functions
- Characterize call sites and arguments to functions
- Malloc tracing
 - Detecting memory leaks
 - **Generating address traces**

■ Error Checking

- C Programming Lab used customized versions of malloc/free to do careful error checking
- Other labs (malloc, shell, proxy) also use interpositioning to enhance checking capabilities

Example program

```
#include <stdio.h>
#include <malloc.h>
#include <stdlib.h>

int main(int argc,
          char *argv[])
{
    int i;
    for (i = 1; i < argc; i++) {
        void *p =
            malloc(atoi(argv[i]));
        free(p);
    }
    return(0);
}
```

int.c

- Goal: trace the addresses and sizes of the allocated and freed blocks, without breaking the program, and without modifying the source code.
- Three solutions: interpose on the library **malloc** and **free** functions at compile time, link time, and load/run time.

Compile-time Interpositioning

```
#ifdef COMPILETIME
#include <stdio.h>
#include <malloc.h>

/* malloc wrapper function */
void *mymalloc(size_t size)
{
    void *ptr = malloc(size);
    printf("malloc(%d)=%p\n", (int)size, ptr);
    return ptr;
}

/* free wrapper function */
void myfree(void *ptr)
{
    free(ptr);
    printf("free(%p)\n", ptr);
}
#endif
```

mymalloc.c

Compile-time Interpositioning

```
#define malloc(size) mymalloc(size)
#define free(ptr) myfree(ptr)

void *mymalloc(size_t size);
void myfree(void *ptr);
```

malloc.h

linux> make intc

```
gcc -Wall -DCOMPILETIME -c mymalloc.c
gcc -Wall -I. -o intc intc.c mymalloc.o
```

linux> make runc

```
./intc 10 100 1000
malloc(10)=0x1ba7010
free(0x1ba7010)
malloc(100)=0x1ba7030
free(0x1ba7030)
malloc(1000)=0x1ba70a0
free(0x1ba70a0)
```

Search for <malloc.h> leads to
/usr/include/malloc.h

Search for <malloc.h> leads to

linux>

Link-time Interpositioning

```
#ifdef LINKTIME
#include <stdio.h>

void * __real_malloc(size_t size);
void __real_free(void *ptr);

/* malloc wrapper function */
void * __wrap_malloc(size_t size)
{
    void *ptr = __real_malloc(size); /* Call libc malloc */
    printf("malloc(%d) = %p\n", (int)size, ptr);
    return ptr;
}

/* free wrapper function */
void __wrap_free(void *ptr)
{
    __real_free(ptr); /* Call libc free */
    printf("free(%p)\n", ptr);
}

#endif
```

mymalloc.c

Link-time Interpositioning

```
linux> make intl
gcc -Wall -DLINKTIME -c mymalloc.c
gcc -Wall -c int.c
gcc -Wall -Wl,--wrap,malloc -Wl,--wrap,free -o intl \
    int.o mymalloc.o

linux> make runl
./intl 10 100 1000
malloc(10) = 0x91a010
free(0x91a010)
. . .
```

Search for <malloc.h> leads to /usr/include/malloc.h

- The “**-Wl**” flag passes argument to linker, replacing each comma with a space.
- The “**--wrap,malloc**” arg instructs linker to resolve references in a special way:
 - Refs to malloc should be resolved as __wrap_malloc
 - Refs to __real_malloc should be resolved as malloc

Load/Run-time Interpositioning

```
#ifdef RUNTIME
#define _GNU_SOURCE
#include <stdio.h>
#include <stdlib.h>
#include <dlfcn.h>

/* malloc wrapper function */
void *malloc(size_t size)
{
    void *(*mallocp)(size_t size);
    char *error;

    mallocp = dlsym(RTLD_NEXT, "malloc"); /* Get addr of libc malloc */
    if ((error = dlerror()) != NULL) {
        fputs(error, stderr);
        exit(1);
    }
    char *ptr = mallocp(size); /* Call libc malloc */
    printf("malloc(%d) = %p\n", (int)size, ptr);
    return ptr;
}
```

Observe that DON'T have
`#include <malloc.h>`

mymalloc.c

Load/Run-time Interpositioning

```
/* free wrapper function */
void free(void *ptr)
{
    void (*freep)(void *) = NULL;
    char *error;

    if (!ptr)
        return;

    freep = dlsym(RTLD_NEXT, "free"); /* Get address of libc free */
    if ((error = dlerror()) != NULL) {
        fputs(error, stderr);
        exit(1);
    }
    freep(ptr); /* Call libc free */
    printf("free(%p)\n", ptr);
}
#endif
```

mymalloc.c

Load/Run-time Interpositioning

```
linux> make intr
gcc -Wall -DRUNTIME -shared -fpic -o mymalloc.so mymalloc.c -ldl
gcc -Wall -o intr int.c
linux> make runr
(LD_PRELOAD="./mymalloc.so" ./intr 10 100 1000)
malloc(10) = 0x91a010
free(0x91a010)
...
linux>
```

Search for <malloc.h> leads to
/usr/include/malloc.h

- The `LD_PRELOAD` environment variable tells the dynamic linker to resolve unresolved refs (e.g., to `malloc`) by looking in `mymalloc.so` first.
- Type into (some) shells as:

```
env LD_PRELOAD=./mymalloc.so ./intr 10 100 1000)
```

Interpositioning Recap

■ Compile Time

- Apparent calls to **malloc/free** get macro-expanded into calls to **mymalloc/myfree**
- Simple approach. Must have access to source & recompile

■ Link Time

- Use linker trick to have special name resolutions
 - **malloc** → **_wrap_malloc**
 - **_real_malloc** → **malloc**

■ Load/Run Time

- Implement custom version of **malloc/free** that use dynamic linking to load library **malloc/free** under different names
- Can use with ANY dynamically linked binary

```
env LD_PRELOAD=./mymalloc.so gcc -c int.c)
```

Linking Recap

- **Usually: Just happens, no big deal**
- **Sometimes: Strange errors**
 - Bad symbol resolution
 - Ordering dependence of linked .o, .a, and .so files
- **For power users:**
 - Interpositioning to trace programs with & without source